

Document Roadmap: Sovereign AI Infrastructure

Bicameral Validator Ladder with Hybrid Logic Router

Project: Local Multi-Model AI System with GPU/CPU Heterogeneous Compute

Version: 1.0

Date: February 5, 2026

Purpose of This Roadmap

This document roadmap guides you through **every document you need to create** from initial concept to operational deployment. As your first project, this roadmap ensures you don't miss critical planning, design, or operational documentation.

Documents are organized in **sequential phases**. Each document builds on previous ones and feeds into the next. Dependencies are clearly marked.

Phase 0: Foundation & Planning Documents

0.1 Technical Analysis Report (COMPLETED)

- **Status:** Already exists (Deep Agent Report.pdf)
- **Purpose:** Comprehensive analysis of the architecture, critical evaluation, feasibility
- **Key Sections:** Architecture components, validation approach, hardware constraints
- **Audience:** Technical team, stakeholders
- **Dependencies:** None
- **Outputs to next phase:** Validated technical approach, identified risks, component specifications

0.2 Comprehensive Architecture Report (COMPLETED)

- **Status:** Already exists (Comprehensive Report - Hybrid Logic Router.pdf)
- **Purpose:** Final iteration architecture with Prolog/Python hybrid, OCR/vision/embeddings
- **Key Sections:** Bicameral design, governance policies, implementation blueprint
- **Audience:** Implementation team
- **Dependencies:** Technical Analysis Report
- **Outputs to next phase:** Detailed architecture, expanded capability requirements

0.3 Product Requirements Document (PRD) ← YOU ARE HERE

- **Status:** In progress
- **Purpose:** Define what you're building, why, for whom, and success criteria
- **Key Sections:** Product vision, user stories, functional/non-functional requirements, success metrics
- **Audience:** All stakeholders, implementation team, future maintainers

- **Dependencies:** Both analysis reports
 - **Outputs to next phase:** Clear requirements baseline, acceptance criteria, scope boundaries
 - **Estimated effort:** 3-5 days
 - **Owner:** Product Lead / Technical Lead
-

Phase 1: Design & Specification Documents

1.1 System Architecture Document (SAD)

- **Purpose:** Translate PRD requirements into detailed technical architecture
- **Key Sections:**
 - Component diagrams (GPU/CPU split, model placement)
 - Data flow diagrams (Generate → Validate → Commit loop)
 - Interface specifications (Python ↔ Prolog, model API contracts)
 - Technology stack details (llama.cpp, SWI-Prolog, vector DB choice)
 - Deployment architecture
- **Audience:** Software engineers, DevOps
- **Dependencies:** PRD, Architecture Reports
- **Outputs to next phase:** Technical blueprints for implementation
- **Estimated effort:** 1-2 weeks
- **Owner:** Solutions Architect / Technical Lead
- **Artifacts:** Architecture diagrams (C4 model recommended), API specs, data models

1.2 Data Architecture & Memory Design Document

- **Purpose:** Define the Markdown memory bus, knowledge graph structure, persistence layer
- **Key Sections:**
 - Markdown file schemas (project_state.md, scratchpad.md, knowledge_graph.md)
 - Vector database schema (embeddings storage)
 - Provenance tracking design (OCR sources, vision input metadata)
 - Backup and versioning strategy
 - Concurrency control mechanisms
- **Audience:** Backend engineers, data engineers
- **Dependencies:** PRD, System Architecture Document
- **Outputs to next phase:** Memory system implementation specs
- **Estimated effort:** 1 week
- **Owner:** Data Architect / Backend Lead

1.3 Routing Logic Specification (Prolog Predicates)

- **Purpose:** Define declarative routing rules, validation policies, grounding requirements
- **Key Sections:**
 - Predicate definitions (routing rules, validation policies, tool selection)
 - Decision tree examples (given input X, why model Y was chosen)
 - Stakes classification logic (low/medium/high)
 - Uncertainty handling rules
 - Grounding policy constraints (OCR/vision source-backing)

- **Audience:** Logic programming developers, system architects
- **Dependencies:** PRD, System Architecture Document
- **Outputs to next phase:** Prolog implementation roadmap
- **Estimated effort:** 1-2 weeks
- **Owner:** Logic Systems Engineer / Routing Lead

1.4 Model Serving & Orchestration Design

- **Purpose:** Define how models are loaded, managed, and orchestrated
- **Key Sections:**
 - Model loading/unloading sequences
 - Warm pool management algorithm
 - GPU ↔ RAM transfer protocols
 - Quantization strategy per model
 - Failure recovery (OOM handling, model swap failures)
 - Monitoring and telemetry points
- **Audience:** ML engineers, DevOps
- **Dependencies:** System Architecture Document
- **Outputs to next phase:** Orchestration implementation specs
- **Estimated effort:** 1 week
- **Owner:** ML Infrastructure Engineer

1.5 Multimodal Pipeline Design Document

- **Purpose:** Specify OCR, vision encoder, and embedding model integration
- **Key Sections:**
 - OCR pipeline (input → extraction → provenance → output)
 - Vision encoder pipeline (image → features → captions → embeddings)
 - Embedding model pipeline (text → vector → retrieval)
 - Tool selection routing (when to invoke each pipeline)
 - Error handling and confidence thresholds
- **Audience:** ML engineers, backend engineers
- **Dependencies:** System Architecture Document, Routing Logic Specification
- **Outputs to next phase:** Multimodal tool implementation specs
- **Estimated effort:** 1 week
- **Owner:** Multimodal Systems Engineer

1.6 Security & Governance Design Document

- **Purpose:** Define security controls, audit mechanisms, governance policies
- **Key Sections:**
 - Threat model (prompt injection, model manipulation, memory tampering)
 - Access controls (who can invoke which models/tools)
 - Audit logging requirements (decision trails, provenance)
 - Data privacy measures (local-only guarantees, encryption at rest)
 - Compliance considerations (if applicable)
- **Audience:** Security engineers, compliance officers
- **Dependencies:** PRD, System Architecture Document
- **Outputs to next phase:** Security implementation requirements

- **Estimated effort:** 1 week
 - **Owner:** Security Architect
-

Phase 2: Implementation Documents

2.1 Development Plan & Sprint Structure

- **Purpose:** Break down implementation into manageable sprints
- **Key Sections:**
 - Sprint breakdown (aligned with phased approach from reports)
 - Team structure and role assignments
 - Development environment setup instructions
 - Code repository structure
 - Branching and merge strategy
 - Definition of Done for each sprint
- **Audience:** Development team, project manager
- **Dependencies:** All Phase 1 design documents
- **Outputs to next phase:** Sprint backlogs, resource allocation
- **Estimated effort:** 3-5 days
- **Owner:** Engineering Manager / Scrum Master

2.2 API Specification Document

- **Purpose:** Define all internal and external API contracts
- **Key Sections:**
 - Python orchestrator API (REST/gRPC endpoints)
 - Prolog query interface (Python ↔ Prolog protocol)
 - Model server APIs (llama.cpp endpoints)
 - Tool APIs (OCR, vision, embeddings)
 - Markdown memory bus access patterns
 - Error response schemas
- **Audience:** Full-stack engineers, integration developers
- **Dependencies:** System Architecture Document
- **Outputs to next phase:** OpenAPI specs, integration test requirements
- **Estimated effort:** 1 week
- **Owner:** API Lead / Backend Architect
- **Artifacts:** OpenAPI/Swagger specs, Postman collections

2.3 Testing Strategy & Test Plan

- **Purpose:** Define comprehensive testing approach
- **Key Sections:**
 - Unit testing strategy (per component)
 - Integration testing approach (router → model → validator flows)
 - End-to-end testing scenarios (50+ representative tasks)
 - Performance testing (latency, throughput, resource utilization)
 - Validation accuracy testing (false positive/negative rates)

- Multimodal pipeline testing (OCR/vision accuracy)
- Load testing (GPU/CPU/RAM stress scenarios)
- Security testing (prompt injection, bypass attempts)
- **Audience:** QA engineers, developers
- **Dependencies:** All Phase 1 design documents
- **Outputs to next phase:** Test cases, test data sets, acceptance criteria
- **Estimated effort:** 1 week
- **Owner:** QA Lead / Test Architect

2.4 Model Prompt Engineering Guide

- **Purpose:** Document system prompts for each model role
- **Key Sections:**
 - Worker prompts (per specialist: coding, reasoning, creative)
 - Validator prompts (governance, structure checking, grounding verification)
 - Router prompts (classification, uncertainty handling)
 - Few-shot examples for each role
 - Prompt versioning and A/B testing strategy
 - Failure mode prompts (retry, clarification, escalation)
- **Audience:** ML engineers, prompt engineers
- **Dependencies:** Routing Logic Specification, Model Serving Design
- **Outputs to next phase:** Production-ready prompts, prompt templates
- **Estimated effort:** 2 weeks (iterative refinement expected)
- **Owner:** Prompt Engineering Lead

Phase 3: Deployment & Operations Documents

3.1 Infrastructure Setup Guide

- **Purpose:** Step-by-step instructions to provision and configure hardware
- **Key Sections:**
 - Hardware procurement checklist (Tesla A2, Xeon, RAM, NVMe)
 - Operating system installation and configuration
 - Driver installation (NVIDIA CUDA, cuDNN)
 - Dependency installation (llama.cpp, SWI-Prolog, Python packages)
 - Filesystem layout (model vault directory structure)
 - Network configuration (if multi-node)
- **Audience:** DevOps engineers, system administrators
- **Dependencies:** System Architecture Document, Model Serving Design
- **Outputs to next phase:** Operational infrastructure
- **Estimated effort:** 3-5 days
- **Owner:** DevOps Lead

3.2 Deployment Runbook

- **Purpose:** Detailed deployment procedures for production
- **Key Sections:**

- Pre-deployment checklist (model downloads, quantization, verification)
- Deployment sequence (services startup order)
- Configuration management (environment variables, config files)
- Health checks and smoke tests
- Rollback procedures
- Blue-green deployment strategy (if applicable)
- **Audience:** DevOps engineers, SREs
- **Dependencies:** Infrastructure Setup Guide
- **Outputs to next phase:** Repeatable deployment process
- **Estimated effort:** 1 week
- **Owner:** DevOps Lead / SRE

3.3 Monitoring & Observability Plan

- **Purpose:** Define what to measure and how to monitor system health
- **Key Sections:**
- Metrics catalog (router accuracy, validation rates, latency, resource utilization)
- Logging strategy (structured logs, audit trails)
- Dashboard designs (Grafana dashboards)
- Alerting rules (OOM warnings, model swap failures, validation bottlenecks)
- Tracing instrumentation (distributed tracing for request flows)
- Performance baselines and SLOs
- **Audience:** SREs, DevOps, operations team
- **Dependencies:** System Architecture Document, Testing Strategy
- **Outputs to next phase:** Monitoring infrastructure, alert definitions
- **Estimated effort:** 1 week
- **Owner:** SRE Lead / Observability Engineer
- **Artifacts:** Prometheus configs, Grafana dashboards, alert rules

3.4 Operations Manual (Runbook)

- **Purpose:** Day-to-day operational procedures
- **Key Sections:**
- Routine maintenance tasks (model updates, log rotation)
- Troubleshooting guide (common failure modes and resolutions)
- Performance tuning procedures
- Backup and recovery procedures
- Incident response protocols
- On-call procedures
- Capacity planning guidelines
- **Audience:** Operations team, on-call engineers
- **Dependencies:** Monitoring Plan, Deployment Runbook
- **Outputs to next phase:** Operational readiness
- **Estimated effort:** 1-2 weeks
- **Owner:** Operations Lead / SRE

3.5 Disaster Recovery & Business Continuity Plan

- **Purpose:** Ensure system can recover from failures

- **Key Sections:**
 - Backup strategy (models, memory files, configuration)
 - Recovery time objectives (RTO) and recovery point objectives (RPO)
 - Failure scenarios and recovery procedures (hardware failure, corruption, etc.)
 - Failover mechanisms (if applicable)
 - Data integrity verification procedures
 - **Audience:** SREs, operations team, management
 - **Dependencies:** System Architecture Document, Operations Manual
 - **Outputs to next phase:** Disaster recovery capability
 - **Estimated effort:** 3-5 days
 - **Owner:** SRE Lead / Operations Manager
-

Phase 4: Knowledge Transfer & Maintenance Documents

4.1 Developer Onboarding Guide

- **Purpose:** Enable new developers to contribute effectively
- **Key Sections:**
- System overview and architecture walkthrough
- Development environment setup
- Code repository navigation
- Key design patterns and conventions
- How to add a new model to the stack
- How to modify routing logic
- How to extend multimodal pipelines
- Debugging techniques and tools
- **Audience:** New developers, contractors
- **Dependencies:** All implementation and operations documents
- **Outputs to next phase:** Self-sufficient development team
- **Estimated effort:** 1 week
- **Owner:** Technical Lead / Senior Engineer

4.2 User Guide / End-User Documentation

- **Purpose:** Teach end users how to interact with the system
- **Key Sections:**
- System capabilities overview
- How to submit requests (interface documentation)
- Understanding routing decisions (transparency features)
- Interpreting validation reports
- Working with OCR/vision inputs
- Best practices for high-quality outputs
- Limitations and known issues
- FAQ and troubleshooting

- **Audience:** End users, domain experts
- **Dependencies:** PRD, all design documents
- **Outputs to next phase:** User adoption
- **Estimated effort:** 1 week
- **Owner:** Technical Writer / Product Manager
- **Artifacts:** User manual (HTML/PDF), video tutorials, quick-start guides

Model Update & Maintenance Guide

- **Purpose:** Document how to update models, prompts, and routing logic
- **Key Sections:**
 - Model evaluation and selection criteria
 - Quantization procedures
 - Model integration checklist
 - Prompt update and testing procedures
 - Routing logic modification guidelines
 - Regression testing requirements
 - Performance impact assessment
- **Audience:** ML engineers, operations team
- **Dependencies:** Model Serving Design, Prompt Engineering Guide
- **Outputs to next phase:** Sustainable maintenance process
- **Estimated effort:** 3-5 days
- **Owner:** ML Lead / Operations Lead

Lessons Learned & Post-Deployment Review

- **Purpose:** Capture insights for future improvements
- **Key Sections:**
 - What worked well
 - What didn't work as expected
 - Performance vs. projections (latency, accuracy, resource usage)
 - Architectural decisions revisited
 - Recommendations for next iteration
 - Technical debt inventory
- **Audience:** All stakeholders, future teams
- **Dependencies:** Operational data, team retrospectives
- **Outputs to next phase:** Improvement roadmap
- **Estimated effort:** 1 week (after 3-6 months operation)
- **Owner:** Technical Lead / Product Manager

Phase 5: Continuous Improvement Documents (Ongoing)

Performance Optimization Reports (Recurring)

- **Purpose:** Document optimization efforts and results
- **Cadence:** Quarterly or after major optimizations

- **Key Sections:** Identified bottlenecks, optimization approach, results, benchmarks

5.2 Incident Post-Mortems (As Needed)

- **Purpose:** Learn from failures and prevent recurrence
- **Trigger:** Any significant outage or issue
- **Key Sections:** Timeline, root cause, impact, remediation, prevention

5.3 Model Performance Evaluations (Recurring)

- **Purpose:** Assess whether current models still meet requirements
- **Cadence:** Quarterly or when new models released
- **Key Sections:** Current performance, new model candidates, evaluation results, recommendation

5.4 Roadmap Updates (Recurring)

- **Purpose:** Evolve the system based on learnings and new requirements
 - **Cadence:** Bi-annually or after major milestones
 - **Key Sections:** Completed features, upcoming features, architectural evolution plans
-

Document Dependencies Visualization

```

Phase 0: Foundation
├─ Technical Analysis Report (✓)
└─ Architecture Report (✓)
└─ PRD (📄 IN PROGRESS)

↓
Phase 1: Design
├─ System Architecture Document ← depends on PRD
├─ Data Architecture ← depends on PRD, SAD
├─ Routing Logic Spec ← depends on PRD, SAD
├─ Model Serving Design ← depends on SAD
├─ Multimodal Pipeline Design ← depends on SAD, Routing Spec
└─ Security & Governance ← depends on PRD, SAD

↓
Phase 2: Implementation
├─ Development Plan ← depends on all Phase 1 docs
├─ API Specification ← depends on SAD
├─ Testing Strategy ← depends on all Phase 1 docs
└─ Prompt Engineering Guide ← depends on Routing Spec, Model Serving

↓
Phase 3: Deployment & Ops
├─ Infrastructure Setup ← depends on SAD, Model Serving
├─ Deployment Runbook ← depends on Infrastructure Setup
├─ Monitoring Plan ← depends on SAD, Testing Strategy
├─ Operations Manual ← depends on Monitoring, Deployment
└─ DR/BC Plan ← depends on SAD, Operations Manual

↓
Phase 4: Knowledge Transfer
├─ Developer Onboarding ← depends on all previous docs
├─ User Guide ← depends on PRD, all design docs
├─ Model Maintenance Guide ← depends on Model Serving, Prompts
└─ Lessons Learned ← depends on operational data

↓
Phase 5: Continuous (ongoing)
├─ Performance Optimization Reports
├─ Incident Post-Mortems
├─ Model Performance Evaluations
└─ Roadmap Updates

```

Document Ownership Matrix

Phase	Document	Primary Owner	Collaborators	Review Required From
0	Technical Analysis	✓ Complete	-	-
0	Architecture Report	✓ Complete	-	-
0	PRD	Product Lead	Tech Lead, Stakeholders	All
1	System Architecture	Solutions Architect	Dev Team	Tech Lead, Security
1	Data Architecture	Data Architect	Backend Lead	Tech Lead, Dev Team
1	Routing Logic Spec	Logic Systems Engineer	Architect	Tech Lead, ML Lead
1	Model Serving Design	ML Infrastructure Eng	DevOps	ML Lead, Architect
1	Multimodal Pipeline	Multimodal Systems Eng	ML Team	ML Lead, Architect
1	Security & Governance	Security Architect	Legal/Compliance	Tech Lead, Management
2	Development Plan	Engineering Manager	Scrum Master	Tech Lead
2	API Specification	API Lead	Backend Team	Architect, Dev Team
2	Testing Strategy	QA Lead	QA Team, Dev Team	Tech Lead, QA Manager
2	Prompt Engineering	Prompt Engineering Lead	ML Team	ML Lead, Tech Lead
3	Infrastructure Setup	DevOps Lead	SRE Team	Tech Lead, Operations
3	Deployment Runbook	DevOps Lead	SRE	Operations Manager
3	Monitoring Plan	SRE Lead	DevOps	

Phase	Document	Primary Owner	Collaborators	Review Required From
				Operations Manager
3	Operations Manual	Operations Lead	SRE Team	Operations Manager
3	DR/BC Plan	SRE Lead	Operations	Operations Manager, Management
4	Developer On-boarding	Technical Lead	Senior Engineers	Engineering Manager
4	User Guide	Technical Writer	Product Manager	Product Lead, UX
4	Model Maintenance	ML Lead	Operations	ML Manager, DevOps
4	Lessons Learned	Tech Lead / PM	All Teams	Management, Stakeholders

Critical Success Factors for Documentation

1. Documentation-Driven Development

- Write design docs BEFORE code
- All major architectural decisions must be documented and reviewed
- PRD is the contract; all work validates against it

2. Version Control

- All documents in Git (Markdown preferred for diff-ability)
- Document version tied to software version
- Change logs maintained for major documents

3. Review Process

- All Phase 1 (Design) docs require architecture review
- All Phase 2 (Implementation) docs require peer review
- All Phase 3 (Operations) docs require ops team sign-off

4. Living Documents

- Documents are never “done”—update as system evolves
- Quarterly review of all Phase 3 (Operations) documents
- Incident learnings feed back into runbooks immediately

5. Accessibility

- Central documentation repository (wiki or Git)
 - Search capability across all documents
 - Clear naming conventions and folder structure
-

Recommended Documentation Tools

For Diagramming:

- **Draw.io / Lucidchart**: Architecture diagrams
- **Mermaid**: Embedded diagrams in Markdown (Git-friendly)
- **C4 Model**: For system architecture (Context, Container, Component, Code)

For Documentation:

- **Markdown + Git**: Version-controlled, diff-friendly, accessible
- **Confluence / Notion**: If team prefers wiki-style
- **Docusaurus / MkDocs**: For static site generation from Markdown

For API Specs:

- **OpenAPI 3.0 / Swagger**: REST API specifications
- **Protobuf**: If using gRPC

For Collaboration:

- **GitHub/GitLab Issues**: Track documentation tasks
 - **Miro / FigJam**: Collaborative whiteboarding
 - **Google Docs**: For draft reviews before committing to Git
-

Next Steps

1. **You are here**: Reviewing Document Roadmap
 2. **Next**: Complete PRD (below this document)
 3. **Then**: Begin Phase 1 (Design Documents), starting with System Architecture Document
 4. **Ongoing**: Maintain document dependencies and update this roadmap as you progress
-

Document Maintenance Schedule

Document Type	Review Frequency	Update Trigger
PRD	Quarterly	Scope changes, major feature requests
Architecture Docs	Bi-annually	Major architectural changes
API Specs	Per release	API changes
Operations Docs	Monthly	Process improvements, incidents
User Guide	Per release	Feature additions, UI changes
Monitoring Plan	Quarterly	New metrics, SLO changes
Security Docs	Quarterly	Threat model changes, audits

Document maintained by: Technical Lead / Product Manager

Last updated: February 5, 2026

Next review: May 5, 2026