

# Data Architecture & Memory Design Document

## Sovereign AI Infrastructure: Transparent Memory Ledger System

**Document Version:** 1.0

**Date:** February 5, 2026

**Status:** Draft for Review

**Owner:** Data Architect / Backend Lead

## Document Control

Version	Date	Author	Changes
1.0	2026-02-05	Data Architecture Team	Initial draft

### Reviewers:

- [ ] Technical Lead
- [ ] Backend Lead
- [ ] ML Lead
- [ ] Security Architect

### Dependencies:

- System Architecture Document (SAD) v1.0
- Product Requirements Document (PRD) v1.0

## Table of Contents

1. Executive Summary
2. Memory Architecture Overview
3. Markdown Memory Ledger
4. Vector Store Architecture
5. Provenance & Audit Trail
6. Concurrency & Consistency
7. Backup & Recovery
8. Access Patterns & Performance
9. Data Lifecycle Management
10. Implementation Specifications
11. Appendices

# 1. Executive Summary

## 1.1 Purpose

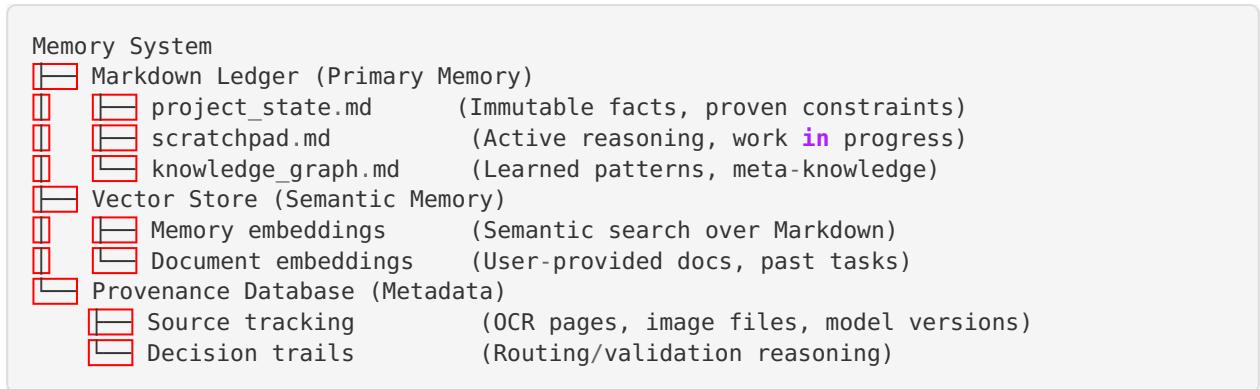
The **Transparent Memory Ledger** is the central nervous system of the Sovereign AI Infrastructure. Unlike traditional databases or neural memory systems, it uses **Markdown files** as a human-readable, auditable, version-controlled shared memory bus that all AI components read from and write to.

## 1.2 Design Philosophy

### Key Principles:

1. **Transparency:** All memory is human-readable Markdown (not binary, not opaque database)
2. **Auditability:** Git version control tracks every change; immutable audit trail
3. **Simplicity:** Plain text files; no complex database schemas; grep-compatible
4. **Provenance:** Every fact includes source attribution (which model, when, why)
5. **Shared Context:** All models (Router, Workers, Validator) access same memory

## 1.3 Core Memory Components



## 1.4 Key Architectural Decisions

Decision	Rationale	Trade-offs
<b>Markdown for primary memory</b>	Human-readable, Git-trackable, LLM-friendly	Not optimized for structured queries vs. SQL
<b>Separate project_state (immutable) vs. scratchpad (mutable)</b>	Clear boundary between proven facts and speculation	More complex state management
<b>Git for versioning</b>	Industry-standard, time-travel capability	Requires Git discipline (commits, etc.)
<b>FAISS for vector store</b>	Fast, CPU-efficient, local-only	Less feature-rich than ChromaDB/Pinecone
<b>File-based (not database)</b>	Simplicity, no database overhead	Manual concurrency control needed

## 2. Memory Architecture Overview

### 2.1 Three-Layer Memory Hierarchy



## 2.2 Memory Flow Diagram

```

graph TD
    Request[User Request] --> Router[Router Classifies]
    Router --> Retrieval[Retrieve Relevant Context<br/>from Memory]
    Retrieval --> Worker[Worker Generates<br/>writes to Scratchpad]
    Worker --> Validator[Validator Checks<br/>reads Scratchpad + Project State]

    Validator -->|PASS| Commit[Commit to Project State<br/>immutable]
    Validator -->|FAIL| Correction[Write Correction<br/>to Scratchpad]
    Correction --> Worker

    Commit --> Extract[Extract Meta-Learnings]
    Extract --> KnowledgeGraph[Update Knowledge Graph]

    Commit --> Embed[Re-embed New Facts]
    KnowledgeGraph --> Embed
    Embed --> VectorStore[Update Vector Store<br/>FAISS]

    style Request fill:#50E3C2,stroke:#333,stroke-width:2px
    style Commit fill:#7ED321,stroke:#333,stroke-width:2px,color:#000
    style VectorStore fill:#4A90E2,stroke:#333,stroke-width:2px,color:#fff

```

## 3. Markdown Memory Ledger

### 3.1 File Structure Overview

**Primary Location:** /opt/sovereign-ai/data/memory/

```

memory/
├── .git/                                # Git repository for version control
├── project_state.md                      # Layer 2: Immutable proven facts
├── scratchpad.md                         # Layer 1: Current reasoning
├── knowledge_graph.md                    # Layer 3: Meta-knowledge
└── sessions/
    ├── 2026-02-01_session.md
    ├── 2026-02-02_session.md
    └── ...
└── README.md                             # Memory system documentation

```

### 3.2 project\_state.md Schema

**Purpose:** Immutable long-term memory of proven, validated facts.

**Structure:**

```

# Project State

--Last Updated--: 2026-02-05 14:30:00 UTC
--Session--: 42
--Git Commit--: abc1234

---

## Global Objectives

--Primary Goal--: [High-level objective, e.g., "Develop payment processing module"]

--Success Criteria--:
- [ ] Criterion 1
- [ ] Criterion 2

---

## Immutable Facts (Validated & Committed)

### Fact #1: [Timestamp: 2026-02-05 10:15:00]
--Actor--: Worker (Qwen Coder 32B)
--Validator--: Granite-H-Small
--Content--:
```python
def calculate_fibonacci(n: int) -> int:
    """Calculate nth Fibonacci number recursively."""
    if n <= 1:
        return n
    return calculate_fibonacci(n-1) + calculate_fibonacci(n-2)
```

```

**Validation:** [PASS] - Syntax ✓, Logic ✓, Type hints ✓

**Source:** User request “Create Fibonacci function”

**Confidence:** 0.95

---

## Fact #2: [Timestamp: 2026-02-05 11:20:00]

**Actor:** Worker (GPT-OSS 20B)

**Validator:** Granite-H-Small

**Content:**

System architecture decision: Use PostgreSQL for transaction storage due to ACID compliance requirements.

**Validation:** [PASS] - Logical consistency ✓, No hallucinations ✓

**Source:** Analysis of payment processing requirements

**Confidence:** 0.89

---

## Proven Constraints

### Hardware Constraints

- GPU VRAM: 16GB (Tesla A2)
- System RAM: 128GB

- Max model size: 32B parameters (Q4 quantization)

## Domain Constraints

- Language: Python 3.10+
  - Framework: FastAPI for REST APIs
  - Database: PostgreSQL 14+
- 

## Dependencies & Environment

**Python Version:** 3.10.12

**Key Libraries:**

- fastapi==0.100.0
  - sqlalchemy==2.0.0
  - psycopg2==2.9.0
- 

## Historical Decisions (Archived)

[Link to past decision logs if needed]

---

**Note:** This file is APPEND-ONLY. Never delete or edit existing facts. To correct, add a new fact superseding the old one.

```

**Key Properties**:
- **Immutability**: Append only; existing entries never modified
- **Timestamped**: Every fact includes UTC timestamp
- **Attributed**: Every fact includes which Actor (Worker/Router/User) and which Validator
- **Validated**: Every fact includes validation verdict and checks performed
- **Searchable**: Plain text; grep-compatible

---

### 3.3 scratchpad.md Schema

**Purpose**: Mutable working memory for current task reasoning.

**Structure**:
```markdown
# Scratchpad (Working Memory)

**Current Task**: [User request summary]
**Started**: 2026-02-05 14:00:00 UTC
**Status**: In Progress
```

---

## Routing Decision

**Router Output**:
- Domain: `coding_architecture`
- Stakes: `high`
- Model: `qwen_coder_32b`
- Validation: `block_by_block`
- Confidence: 0.92
- Reasoning: "Complex multi-file refactoring; architectural focus indicated by keywords: design patterns, SOLID principles"

---

## Active Reasoning Stream

### Step 1: Architecture Design [PENDING]
**Actor**: Worker (Qwen Coder 32B)

**Context**:
- User request: "Refactor the payment module to follow SOLID principles"
- Current state: Monolithic PaymentService class (500 lines)
- Goal: Separate into Single Responsibility components

**Generated Output**:
```markdown
Proposed Architecture:
1. PaymentProcessor (core logic)
2. PaymentValidator (validation rules)
3. PaymentLogger (audit trail)
4. PaymentNotifier (external notifications)
```

```

**Status:** Awaiting validation

---

## Validator Feedback on Step 1 [FAIL]

**Validator:** Granite-H-Small

**Timestamp:** 2026-02-05 14:05:00 UTC

**Verdict:** [FAIL]

**Reasoning:**

The proposed architecture is missing:

1. PaymentRepository (data access layer) - violates dependency inversion
2. Error handling strategy not specified
3. Interface definitions (abstractions) not provided

**Correction Directive:**

Add PaymentRepository for data access, define interfaces (IPaymentProcessor, etc.), and specify error handling approach (exceptions vs. result types).

---

## Step 1: Architecture Design [RETRY #1]

**Actor:** Worker (Qwen Coder 32B)

**Timestamp:** 2026-02-05 14:07:00 UTC

**Revised Output:**

Proposed Architecture:

1. IPaymentProcessor (interface)
  - Concrete: StripePaymentProcessor, PayPalPaymentProcessor
2. IPaymentValidator (interface)
  - Concrete: SchemaValidator, BusinessRuleValidator
3. IPaymentRepository (interface)
  - Concrete: PostgresPaymentRepository
4. PaymentLogger (cross-cutting concern)
5. PaymentNotifier (event-driven notifications)

Error Handling: Use Result<T, Error> pattern (no exceptions for business logic)

**Status:** Awaiting validation

---

## Validator Feedback on Step 1 Retry #1 [PASS]

**Validator:** Granite-H-Small

**Timestamp:** 2026-02-05 14:10:00 UTC

**Verdict:** [PASS]

**Reasoning:**

- Interfaces defined ✓
- Dependency inversion respected ✓
- Repository pattern included ✓
- Error handling specified ✓

**Action:** Committing architecture design to project\_state.md

---

## Temporary Notes (Scratchable)

---

- User mentioned they prefer explicit error handling over exceptions
  - Current codebase uses pytest for testing; maintain compatibility
  - Consider adding PaymentMetrics for observability (deferred to later step)
- 

**Scratchpad cleared after task completion or session end.**

```
**Key Properties**:
- **Mutability**: High; frequently updated, cleared after task completion
- **Feedback Loop**: Contains validator corrections, retry attempts
- **Temporary**: Not intended for long-term storage (archived after session)
- **Append-heavy**: New steps appended sequentially
```

---

### **### 3.4 knowledge\_graph.md Schema**

**Purpose**: Meta-knowledge **and** learned patterns about models, domains, best practices.

**Structure**:

```
```markdown
# Knowledge Graph (Meta-Knowledge)
```

**Last Updated**: 2026-02-05 14:30:00 UTC

---

## **## Model-Specific Learnings**

### **### Qwen Coder 32B**

**Strengths**:

- Excellent at multi-file refactoring
- Strong architectural reasoning
- Accurate **type** inference

**Weaknesses**:

- Requires explicit **type** hints **for** validation to pass
- Occasionally over-engineers simple solutions

**Best Practices**:

- Provide full file context (**not** just snippets)
- Explicitly request "**pragmatic**" solutions **if** simplicity desired

**Last Updated**: 2026-02-05

---

### **### Nemotron 30B**

**Strengths**:

- Fast, practical implementations
- Performance-oriented code

**Weaknesses**:

- Hallucinates Rust **library** APIs (low-confidence on Rust)
- Sometimes skips edge **case** handling

**Best Practices**:

- Use **for** Python, C++, Go implementations
- Avoid **for** Rust unless you verify API existence
- Explicitly request edge **case** handling

**Last Updated**: 2026-02-04

---

### **## GPT-OSS 20B**

**Strengths**:

- Balanced reasoning

- Good at planning **and** decomposition

**\*\*Weaknesses\*\*:**

- Less specialized than coding models
- Verbose explanations (can be edited down)

**\*\*Best Practices\*\*:**

- Use **for** general reasoning, **not** code generation
- Good **for** architectural planning before implementation

**\*\*Last Updated\*\*:** 2026-02-03

---

### **### Granite-H-Small (Validator)**

**\*\*Strengths\*\*:**

- Strict validation, low **false** negative rate
- Good at catching hallucinations

**\*\*Weaknesses\*\*:**

- Can be overly strict (**false** positives on creative solutions)
- Validation prompts need tuning **for** domain

**\*\*Best Practices\*\*:**

- For creative domains, set validation to "**lenient**" mode
- Provide clear validation criteria upfront

**\*\*Last Updated\*\*:** 2026-02-05

---

### **## Domain-Specific Patterns**

#### **### Python Coding**

**\*\*Common Issues\*\*:**

- Forgetting **type** hints → validation fails
- Mixing sync/async without explicit markers

**\*\*Best Practices\*\*:**

- Always request **type**-hinted code
- Specify sync vs. async upfront

---

#### **### Architecture Design**

**\*\*Common Issues\*\*:**

- Missing **interface** definitions
- Not considering testability

**\*\*Best Practices\*\*:**

- Always define interfaces (abstractions)
- Request test strategy alongside architecture

---

### **## Validation Patterns**

#### **### High False Positive Scenarios**

- Creative writing (validator too literal)
- Experimental code (validator expects production-ready)

**\*\*Mitigation\*\*:** Set validation policy to "**lenient**" **or** "**end\_stage**" instead of "**block\_by\_block**"

```

---  

### High False Negative Scenarios  

- Subtle logic errors in complex algorithms  

- Hallucinated APIs in less-common languages  

**Mitigation**: Multiple validation passes; cross-check with external docs  

---  

## User Preferences (Session-Specific)  

**Current User**:  

- Prefers explicit error handling (Result types) over exceptions  

- Values simplicity over clever solutions  

- Uses pytest for testing  

---  

**Note**: This file evolves over time as patterns emerge. Review quarterly.

```

#### **Key Properties:**

- **Meta-knowledge:** Not facts about the project, but about the system itself
  - **Evolving:** Updated as patterns are discovered
  - **Actionable:** Directly informs routing and validation decisions
  - **Prunable:** Old/obsolete learnings can be archived
- 

## 3.5 Markdown Parsing & Generation

#### **Parser Requirements:**

- Parse Markdown headings (H1-H6) for structure
- Extract code blocks (`language ...`)
- Extract lists, tables, links
- Preserve formatting

**Python Implementation** (using `markdown-it-py`):

```

from markdown_it import MarkdownIt
from pathlib import Path

class MemoryParser:
    def __init__(self, memory_dir: Path):
        self.memory_dir = memory_dir
        self.md = MarkdownIt()

    def parse_project_state(self) -> Dict[str, Any]:
        """Parse project_state.md into structured data."""
        filepath = self.memory_dir / "project_state.md"
        content = filepath.read_text()

        # Parse Markdown
        tokens = self.md.parse(content)

        # Extract sections
        facts = self._extract_facts(tokens)
        constraints = self._extract_constraints(tokens)
        objectives = self._extract_objectives(tokens)

        return {
            "facts": facts,
            "constraints": constraints,
            "objectives": objectives,
            "raw_content": content
        }

    def _extract_facts(self, tokens) -> List[Dict]:
        """Extract immutable facts from parsed tokens."""
        facts = []
        # Implementation: Find sections starting with "### Fact #"
        # Extract timestamp, actor, validator, content
        return facts

```

## 4. Vector Store Architecture

### 4.1 Purpose & Design

#### Why Vector Store?

- **Semantic Retrieval:** Find relevant context even if keywords don't match
- **Scalability:** Memory ledger grows over time; text search alone becomes slow
- **Context Augmentation:** Provide Workers with relevant past facts without full memory dump

#### Technology Choice: FAISS (Facebook AI Similarity Search)

- CPU-efficient (no GPU needed)
- Local-only (no external dependencies)
- Fast approximate nearest neighbor search
- Lightweight (no database server required)

### 4.2 Vector Store Schema

**Storage Location:** /opt/sovereign-ai/data/vector\_store/

```

vector_store/
└── memory_embeddings.index      # FAISS index for memory entries
└── memory_metadata.json         # Metadata (timestamps, sources, etc.)
└── document_embeddings.index    # FAISS index for user docs (OCR outputs, etc.)
└── document_metadata.json       # Metadata for documents

```

#### **Embedding Model:** sentence-transformers/all-MiniLM-L6-v2

- Dimension: 384
- Speed: Fast (CPU-efficient)
- Quality: Good for semantic search

## 4.3 Embedding Strategy

#### **What Gets Embedded:**

Source	Content to Embed	Update Frequency
<b>project_state.md</b>	Each fact (as individual chunks)	After each commit
<b>knowledge_graph.md</b>	Each model/domain section	Daily or after updates
<b>scratchpad.md</b>	NOT embedded (temporary, high churn)	N/A
<b>User documents</b>	OCR-extracted text (paragraphs)	On document ingestion

#### **Chunking Strategy:**

- **Chunk size:** 200-300 words (~3-4 sentences)
- **Overlap:** 50 words between consecutive chunks (for context continuity)
- **Metadata:** Each chunk tagged with source (file, section, timestamp)

#### **Example Embedding Entry:**

```
{
  "id": "fact_42_20260205_1015",
  "embedding": [0.123, -0.456, 0.789, ...],  # 384-dim vector
  "metadata": {
    "source_file": "project_state.md",
    "section": "Fact #42",
    "timestamp": "2026-02-05T10:15:00Z",
    "actor": "qwen_coder_32b",
    "content_preview": "def calculate_fibonacci(n: int) -> int: ..."
  }
}
```

## 4.4 Retrieval Process

#### **Query Flow:**

```
User Query → Embed Query → FAISS Search (top-k) → Fetch Metadata → Return Contexts
```

### Python Implementation:

```

import faiss
import numpy as np
from sentence_transformers import SentenceTransformer

class VectorMemory:
    def __init__(self, index_path: str):
        self.index = faiss.read_index(index_path)
        self.embedder = SentenceTransformer('all-MiniLM-L6-v2')
        self.metadata = self._load_metadata()

    def retrieve_relevant(self, query: str, top_k: int = 5) -> List[Dict]:
        """Retrieve top-k relevant memory entries."""
        # Embed query
        query_vector = self.embedder.encode([query])[0]

        # Search FAISS index
        distances, indices = self.index.search(
            np.array([query_vector], dtype=np.float32),
            top_k
        )

        # Fetch metadata
        results = []
        for idx, dist in zip(indices[0], distances[0]):
            results.append({
                "content": self.metadata[idx]["content_preview"],
                "source": self.metadata[idx]["source_file"],
                "timestamp": self.metadata[idx]["timestamp"],
                "relevance_score": 1.0 / (1.0 + dist)  # Convert distance to similarity
            })

        return results

```

### Performance Target:

- Retrieval latency: <100ms for top-10 results
- Index size: ~10MB per 10,000 embeddings

## 5. Provenance & Audit Trail

### 5.1 Provenance Tracking

**Definition:** Provenance = “where did this data come from?”

#### Critical for:

- **Compliance:** Regulators need to trace AI decisions back to sources
- **Debugging:** Understand why output was generated
- **Grounding:** Ensure claims are backed by evidence (OCR text, images)

#### Provenance Metadata Structure:

```

@dataclass
class ProvenanceMetadata:
    """Tracks origin of a fact or output."""

    # Who created this?
    actor: str          # "qwen_coder_32b", "user", "ocr_pipeline"
    actor_version: str  # Model version or tool version

    # When?
    timestamp: datetime # UTC

    # Why/How?
    trigger: str         # "user_request", "validation_correction",
    "auto_learning"
    parent_request_id: str # Link back to originating user request

    # Source materials
    source_documents: List[str] # ["scanned_doc_page_3.png", "user_prompt.txt"]
    source_citations: List[str] # ["OCR line 42-45", "Figure 2 caption"]

    # Validation trail
    validator: Optional[str]   # "granite_h_small" or None
    validation_verdict: str     # "PASS", "FAIL", "NOT_VALIDATED"

    # Confidence
    confidence_score: float    # 0.0-1.0
    uncertainty_notes: str      # E.g., "OCR unclear on this section"

```

## 5.2 Audit Trail in project\_state.md

Every fact in `project_state.md` includes inline provenance:

```

### Fact #42: [Timestamp: 2026-02-05 10:15:00]
**Actor**: Worker (Qwen Coder 32B v1.5.0)
**Validator**: Granite-H-Small v2.0.1
**Source**: User request "Create Fibonacci function"
**Parent Request**: req_1234567890
**Validation**: [PASS] - Syntax ✓, Logic ✓, Type hints ✓
**Confidence**: 0.95
**Provenance**:
- Trigger: User request
- No source documents (generated from scratch)
- Validation checks: syntax_check (PASS), logic_check (PASS), hallucination_check (PASS)

**Content**:
[... actual fact content ...]

```

## 5.3 Multimodal Provenance

For OCR-extracted text:

### ### Fact #55: Patient Diagnosis

**\*\*Actor\*\*:** OCR Pipeline (Tesseract 5.0)  
**\*\*Source Document\*\*:** medical\_record\_page\_3.png

**\*\*OCR Confidence\*\*:** 0.94

**\*\*Extracted Text\*\*:**

"Patient diagnosed with Type 2 Diabetes on 2025-08-15."

**\*\*Provenance\*\*:**

- Source: medical\_record\_page\_3.png, bounding box (120, 450, 680, 490)
- OCR engine: Tesseract 5.0.3
- Preprocessing: Grayscale conversion, noise reduction
- Uncertain regions: Date ("2025-08-15" confidence 0.89, possible "2025-08-18")

**\*\*Grounding\*\*:** All subsequent claims about this diagnosis MUST cite this fact.

### For Vision-encoded images:

### ### Fact #67: Chart Analysis

**\*\*Actor\*\*:** Vision Encoder (CLIP ViT-B/32)

**\*\*Source Document\*\*:** sales\_chart\_q4.png

**\*\*Generated Caption\*\*:**

"Bar chart showing quarterly sales: Q1 \$2.3M, Q2 \$2.8M, Q3 \$3.1M, Q4 \$3.5M"

**\*\*Provenance\*\*:**

- Source: sales\_chart\_q4.png (sha256: abc123...)
- Vision model: CLIP ViT-B/32
- Confidence: Q1-Q3 values (0.92), Q4 value (0.78 - partially occluded)
- Objects detected: 4 bars, x-axis labels, y-axis scale

**\*\*Grounding\*\*:** Any claims about Q4 sales must acknowledge lower confidence (0.78).

## 5.4 Audit Trail Export

**Format:** PDF or HTML report

**Contents:**

- Complete timeline (routing → generation → validation → commit)
- All provenance metadata
- Source documents (embedded images, OCR text)
- Decision reasoning (why this model, why this validation policy)

**Use Case:** Compliance audits, debugging, user transparency

## 6. Concurrency & Consistency

### 6.1 Concurrency Challenges

**Problem:** Multiple processes/threads accessing Markdown files simultaneously:

- **Reader-Writer conflict:** Validator reading scratchpad while Worker writing
- **Write-Write conflict:** Two writers attempting simultaneous append (rare in v1.0 single-task execution)

### 6.2 Locking Strategy

**File-Level Locking** (using Python `filelock`):

```

from filelock import FileLock
from pathlib import Path

class MemoryManager:
    def __init__(self, memory_dir: Path):
        self.memory_dir = memory_dir
        self.locks = {
            "project_state": FileLock(memory_dir / "project_state.lock"),
            "scratchpad": FileLock(memory_dir / "scratchpad.lock"),
            "knowledge_graph": FileLock(memory_dir / "knowledge_graph.lock"),
        }

    def append_to_scratchpad(self, entry: str):
        """Atomic append to scratchpad with locking."""
        with self.locks["scratchpad"]:
            filepath = self.memory_dir / "scratchpad.md"
            with open(filepath, 'a', encoding='utf-8') as f:
                f.write(f"\n{entry}\n")
                f.flush() # Ensure write to disk

    def commit_to_project_state(self, fact: str):
        """Atomic commit to project_state with locking."""
        with self.locks["project_state"]:
            filepath = self.memory_dir / "project_state.md"
            with open(filepath, 'a', encoding='utf-8') as f:
                f.write(f"\n---\n{fact}\n")
                f.flush()

```

#### **Locking Rules:**

- **Shared reads:** Multiple readers allowed (no lock needed)
- **Exclusive writes:** Only one writer at a time (acquire lock)
- **Timeout:** 10 seconds (if lock not acquired, fail gracefully)

## **6.3 Consistency Guarantees**

#### **Atomicity:**

- Each append operation is atomic (file lock ensures no partial writes)
- Git commits provide transaction-like behavior (all-or-nothing)

#### **Durability:**

- All writes flushed to disk immediately (`f.flush()`)
- Git commits create permanent snapshots

#### **Consistency:**

- Schema validation on write (ensure Markdown structure preserved)
- No orphaned references (e.g., scratchpad referencing non-existent fact)

#### **Isolation** (weak in v1.0):

- Single-task execution means minimal concurrency
- Future: Implement MVCC (Multi-Version Concurrency Control) if multi-task support added

## **6.4 Conflict Resolution**

**Scenario:** Validator and Worker both try to write to scratchpad simultaneously (rare)

#### **Resolution:**

1. First writer acquires lock, completes write

2. Second writer blocks until lock released
3. Second writer then appends (sequential, not lost)

**Future Enhancement:** Operational Transform (OT) or CRDTs for true concurrent editing

---

## 7. Backup & Recovery

### 7.1 Backup Strategy

#### Frequency:

- **Continuous:** Git commits after every task completion
- **Daily:** Full snapshot of memory directory
- **Weekly:** Full snapshot of model vault + memory (off-site if possible)

**Backup Script** (example):

```
#!/bin/bash
# daily_backup.sh

DATE=$(date +%Y-%m-%d)
MEMORY_DIR="/opt/sovereign-ai/data/memory"
BACKUP_DIR="/mnt/backups/memory"

# Create backup directory
mkdir -p "$BACKUP_DIR/$DATE"

# Copy memory ledger
cp -r "$MEMORY_DIR" "$BACKUP_DIR/$DATE/"

# Git bundle (complete history)
cd "$MEMORY_DIR"
git bundle create "$BACKUP_DIR/$DATE/memory.bundle" --all

# Compress
cd "$BACKUP_DIR"
tar -czf "$DATE.tar.gz" "$DATE/"
rm -rf "$DATE/" # Remove uncompressed

# Cleanup old backups (keep 30 days)
find "$BACKUP_DIR" -name "*.tar.gz" -mtime +30 -delete

echo "Backup completed: $BACKUP_DIR/$DATE.tar.gz"
```

#### Backup Verification:

- **Checksum:** SHA256 hash of backup archives
- **Test Restore:** Monthly test restore to verify backup integrity

### 7.2 Disaster Recovery

#### Failure Scenarios:

Scenario	Impact	Recovery Procedure	RTO	RPO
<b>Memory file corruption</b>	Data loss, system unusable	Restore from Git history or daily backup	5 min	Last commit (~1 task)
<b>Disk failure (NVMe)</b>	Complete data loss	Restore from weekly backup	1 hour	Up to 7 days of data
<b>Accidental deletion</b>	Specific facts lost	Git revert or restore from backup	2 min	0 (Git history)
<b>Git repository corruption</b>	Version control broken	Re-clone from backup bundle	10 min	Last daily backup

#### Recovery Process (Memory File Corruption):

```

# 1. Detect corruption (checksum mismatch)
# 2. Stop orchestrator
systemctl stop sovereign-orchestrator

# 3. Move corrupted file
mv /opt/sovereign-ai/data/memory/project_state.md /tmp/corrupted_project_state.md

# 4. Restore from Git
cd /opt/sovereign-ai/data/memory
git checkout HEAD -- project_state.md

# 5. Verify integrity
sha256sum project_state.md # Compare with known good checksum

# 6. Restart orchestrator
systemctl start sovereign-orchestrator

```

### 7.3 Point-in-Time Recovery (Git)

**Use Case:** Rollback to state from 3 days ago

```

cd /opt/sovereign-ai/data/memory

# 1. Find commit from 3 days ago
git log --since="3 days ago" --until="2 days ago" --oneline

# 2. Create new branch from that commit (non-destructive)
git checkout -b recovery_branch abc1234

# 3. Inspect state, verify
cat project_state.md

# 4. If correct, merge or replace main branch
git checkout main
git reset --hard abc1234

# 5. Resume operations

```

## 8. Access Patterns & Performance

### 8.1 Common Access Patterns

Operation	Frequency	Latency Target	Implementation
<b>Read project_state</b>	Every validation (high)	<50ms	In-memory caching (TTL 60s)
<b>Append to scratch-pad</b>	Every Worker generation	<20ms	Direct file append (buffered)
<b>Commit to project_state</b>	Every validation PASS	<30ms	File append + Git commit (async)
<b>Semantic retrieval</b>	Every task start	<100ms	FAISS search
<b>Full memory scan</b>	Rare (manual debugging)	<1s	Full file read + parse

### 8.2 Caching Strategy

**In-Memory Cache** (Python `functools.lru_cache`):

```

import functools
import time
from typing import Dict, Any

class MemoryManager:
    def __init__(self):
        self._cache_timestamp = {}
        self._cache_ttl = 60 # seconds

    @functools.lru_cache(maxsize=1)
    def read_project_state_cached(self) -> Dict[str, Any]:
        """Cached read of project_state with TTL."""
        cache_key = "project_state"
        now = time.time()

        # Invalidate cache if TTL expired
        if cache_key in self._cache_timestamp:
            if now - self._cache_timestamp[cache_key] > self._cache_ttl:
                self.read_project_state_cached.cache_clear()

        # Update timestamp
        self._cache_timestamp[cache_key] = now

        # Read (will be cached by lru_cache)
        return self._parse_project_state()

```

#### **Cache Invalidation:**

- **Time-based:** TTL of 60 seconds
- **Event-based:** Clear cache after commit to project\_state
- **Manual:** CLI command `sovereign cache clear`

## **8.3 Performance Optimization**

### **Optimization 1: Incremental Git Commits**

- Problem: Git commit after every fact is slow (100-200ms)
- Solution: Batch commits every 5 facts or every 5 minutes

### **Optimization 2: Lazy Parsing**

- Problem: Parsing entire project\_state.md is slow as it grows
- Solution: Parse only relevant sections (e.g., last N facts, specific section)

### **Optimization 3: Scratchpad Pruning**

- Problem: Scratchpad grows unbounded during long tasks
- Solution: Prune validated/committed blocks from scratchpad (move to project\_state)

### **Optimization 4: FAISS Index Updates**

- Problem: Rebuilding entire FAISS index after every commit is slow
- Solution: Incremental index updates (add new embeddings without rebuild)

## **8.4 Scalability Projections**

#### **Memory Ledger Growth:**

- Assumption: 100 tasks/day, 10KB per task
- Growth: 1MB/day, ~365MB/year
- Manageable for years without performance degradation

**Vector Store Growth:**

- Assumption: 100 embeddings/day, 384 dims × 4 bytes = 1.5KB per embedding
- Growth: 150KB/day, ~54MB/year (index size)
- FAISS efficient for millions of vectors

**Git Repository Size:**

- Text-heavy (Markdown), compresses well
- Expected: <100MB per year (with history)

## 9. Data Lifecycle Management

### 9.1 Lifecycle Stages



### 9.2 Data Retention Policies

Data Type	Retention Period	Archive After	Delete After	Rationale
<code>project_state.md</code>	Indefinite	Never	Never	Immutable facts; compliance
<code>scratchpad.md</code>	Session	End of session	After archiving	Temporary working memory
<code>knowledge_graph.md</code>	Indefinite	Annually (old patterns)	Manually (if obsolete)	Meta-knowledge evolves slowly
<b>Logs</b>	30 days active	30 days	90 days	Compliance, debugging
<b>Backups</b>	30 days (daily), 1 year (weekly)	Immediate	After retention	Disaster recovery

### 9.3 Archiving Process

**When:** End of session, daily, or manually

**Process:**

```

def archive_session():
    """Archive current session scratchpad to sessions/ directory."""
    from datetime import datetime

    date_str = datetime.utcnow().strftime("%Y-%m-%d")
    session_file = f"sessions/{date_str}_session.md"

    # Copy scratchpad to archive
    shutil.copy("scratchpad.md", session_file)

    # Clear scratchpad (or reset to template)
    with open("scratchpad.md", 'w') as f:
        f.write("# Scratchpad (Working Memory)\n\n[Empty - New Session]\n")

    # Git commit
    subprocess.run(["git", "add", session_file, "scratchpad.md"])
    subprocess.run(["git", "commit", "-m", f"Archive session {date_str}"])

    print(f"Session archived: {session_file}")

```

## 9.4 Data Deletion (GDPR Compliance)

**Use Case:** User requests data deletion (GDPR “right to be forgotten”)

**Challenge:** project\_state.md is immutable; how to delete?

**Solution:**

1. **Redaction:** Replace sensitive content with [REDACTED per GDPR request]
2. **Git History Rewrite:** Use git filter-branch or BFG Repo-Cleaner to remove from history
3. **New Commit:** Commit redacted version
4. **Backup Purge:** Delete old backups containing sensitive data

**Example:**

```

# Redact specific fact from project_state.md
sed -i 's/John Doe/[REDACTED]/g' project_state.md

# Rewrite Git history to remove all instances
git filter-branch --tree-filter 'sed -i "s/John Doe/[REDACTED]/g" project_state.md' HEAD

# Force push (if remote)
git push --force

```

**Note:** This is complex and should be rare. Design data collection to minimize PII.

## 10. Implementation Specifications

### 10.1 Python MemoryManager Class

**Interface:**

```

from pathlib import Path
from typing import Dict, List, Any
from dataclasses import dataclass
from datetime import datetime

@dataclass
class MemoryEntry:
    timestamp: datetime
    actor: str
    entry_type: str # "fact", "reasoning", "correction", "learning"
    content: str
    metadata: Dict[str, Any]

class MemoryManager:
    """
    Central manager for Markdown memory ledger and vector store.
    """

    def __init__(self, memory_dir: Path):
        self.memory_dir = memory_dir
        self.project_state_path = memory_dir / "project_state.md"
        self.scratchpad_path = memory_dir / "scratchpad.md"
        self.knowledge_graph_path = memory_dir / "knowledge_graph.md"
        self.vector_store = VectorMemory(memory_dir / "../vector_store")

    # --- READ OPERATIONS ---

    def read_project_state(self) -> Dict[str, Any]:
        """Read and parse project_state.md"""
        pass

    def read_scratchpad(self) -> str:
        """Read current scratchpad"""
        pass

    def read_knowledge_graph(self) -> Dict[str, Any]:
        """Read and parse knowledge_graph.md"""
        pass

    def get_relevant_context(self, query: str, top_k: int = 5) -> List[str]:
        """Semantic retrieval from vector store"""
        return self.vector_store.retrieve_relevant(query, top_k)

    # --- WRITE OPERATIONS ---

    def append_to_scratchpad(self, entry: MemoryEntry) -> bool:
        """Atomic append to scratchpad"""
        pass

    def commit_to_project_state(self, fact: MemoryEntry) -> bool:
        """Commit validated fact to project_state (immutable)"""
        pass

    def update_knowledge_graph(self, learning: str) -> bool:
        """Add new learning to knowledge_graph"""
        pass

    # --- UTILITY OPERATIONS ---

    def clear_scratchpad(self):
        """Clear scratchpad (start new session)"""
        pass

```

```

def archive_session(self, session_name: str):
    """Archive current scratchpad to sessions"""
    pass

def git_commit(self, message: str):
    """Commit current state to Git"""
    pass

def verify_integrity(self) -> bool:
    """Check for corruption, validate schemas"""
    pass

```

## 10.2 Configuration

**YAML Configuration** ( config/memory.yaml ):

```

memory:
  base_dir: "/opt/sovereign-ai/data/memory"

  files:
    project_state: "project_state.md"
    scratchpad: "scratchpad.md"
    knowledge_graph: "knowledge_graph.md"

  cache:
    enabled: true
    ttl_seconds: 60

  git:
    auto_commit: true
    commit_batch_size: 5 # Commit after every 5 facts
    commit_interval_minutes: 5 # Or every 5 minutes

  backup:
    enabled: true
    daily_backup_time: "02:00" # 2 AM
    retention_days: 30
    backup_dir: "/mnt/backups/memory"

  locking:
    timeout_seconds: 10

vector_store:
  base_dir: "/opt/sovereign-ai/data/vector_store"
  embedding_model: "sentence-transformers/all-MiniLM-L6-v2"
  index_type: "FAISS_FLAT" # or FAISS_IVF for larger datasets

chunking:
  chunk_size_words: 250
  overlap_words: 50

retrieval:
  default_top_k: 5
  relevance_threshold: 0.7 # Min similarity score

```

## 10.3 Testing Strategy

**Unit Tests:**

- test\_memory\_manager.py : Test read/write operations, locking, parsing

- `test_vector_store.py` : Test embedding, retrieval, index updates
- `test_provenance.py` : Test metadata tracking, audit trail generation

#### **Integration Tests:**

- `test_end_to_end_memory.py` : Full flow: generate → validate → commit → retrieve

#### **Performance Tests:**

- `test_memory_performance.py` : Measure read/write latencies, cache effectiveness

#### **Corruption Tests:**

- `test_recovery.py` : Simulate file corruption, test recovery from Git/backup
- 

## **11. Appendices**

### **11.1 Markdown Formatting Best Practices**

#### **Consistent Structure:**

- Use H1 (#) for file title
- Use H2 (##) for major sections
- Use H3 (###) for sub-sections (e.g., individual facts)
- Use H4+ for nested content

#### **Code Blocks:**

- Always specify language: ````python`, not `````
- Use triple backticks, not indentation

#### **Lists:**

- Use - for unordered lists
- Use 1. , 2. for ordered lists

#### **Tables:**

- Use Markdown tables for structured data
- Always include header row

#### **Links:**

- Use `[text](url)` format
- Internal links: `[See Fact #42] (#fact-42)`

### **11.2 Schema Evolution**

#### **Versioning:**

- Add schema version to top of each memory file
- Example: `<!-- Schema Version: 1.0 -->`

#### **Backward Compatibility:**

- New fields are optional (don't break old parsers)
- Deprecation period: 6 months before removing old fields

#### **Migration Process:**

1. Add new schema version to parser
2. Support both old and new (dual-read)
3. Background migration (convert old → new)
4. After 6 months, drop old schema support

## 11.3 Debugging & Troubleshooting

### Common Issues:

Issue	Symptom	Diagnosis	Fix
<b>Slow reads</b>	High latency on project_state reads	File too large (>10MB)	Prune old facts, move to archive
<b>Lock timeout</b>	Write operations fail	Multiple writers, lock held too long	Reduce lock hold time, check for deadlocks
<b>Git merge conflicts</b>	Git commit fails	Concurrent commits (shouldn't happen in v1.0)	Manual merge, enforce sequential execution
<b>Vector search poor quality</b>	Irrelevant results	Embedding model mismatch, outdated index	Re-index with correct model
<b>Corruption detected</b>	Checksum mismatch	Disk error, partial write	Restore from Git or backup

### Diagnostic Commands:

```
# Check memory file integrity
sovereign memory verify

# View recent memory commits
sovereign memory log --last 10

# Force cache clear
sovereign memory cache clear

# Manual Git status
cd /opt/sovereign-ai/data/memory && git status

# Re-index vector store
sovereign memory reindex --force
```

**Document Status:** Draft for Review

**Next Steps:** Review by Technical Lead, Backend Lead, ML Lead, Security Architect

**Target Approval Date:** 2026-02-12

**Owner:** Data Architect / Backend Lead