

### 3. The Router Specification (Skeleton)

Language: SWI-Prolog

Integration: Python pyswip wrapper

#### 3.1 Term Shapes

The following Prolog predicates define the core routing logic interface.

- `route(+Request, -Decision)`: Main entry point.
- `classify_domain(+Request, -Domain)`: Determines domain category.
- `assess_stakes(+Request, +Domain, -Stakes)`: Calculates stakes level.
- `select_model(+Domain, +Stakes, -Model, -Confidence)`: Picks best worker.
- `validation_policy(+Stakes, -Policy)`: Determines validation rigor.
- `detect_tools(+Request, -ToolsList)`: Identifies multimodal needs.
- `overall_confidence(+Domain, +Stakes, +Model, +ModelConf, -TotalConf)`: Aggregates confidence.

#### 3.2 Router Input Contract

The `route/2` predicate accepts a Request atom or string.

- Input: Request (String) - The raw user prompt (normalized by Python before passing).
- Context: (Optional) Passed as auxiliary facts if needed (e.g., `user_preference(speed)`).

#### 3.3 Router Output Contract

The Decision variable returns a Dict or JSON-compatible structure:

prolog

Copy

```
Decision = {  
    domain: Atom,      % e.g., coding_architecture  
    stakes: Atom,      % e.g., high  
    recommended_model: Atom, % e.g., qwen_coder_32b  
    validation_policy: Atom, % e.g., block_by_block  
    tools_required: List,  % e.g., [ocr, embeddings]  
    confidence: Float,   % 0.0 - 1.0  
    reasoning: String    % Explanation of the decision  
}
```

#### 3.4 Error Semantics

- Failure to Match: If `classify_domain` fails, the system defaults to unknown domain.

- Low Confidence: If overall\_confidence < 0.75, the uncertainty flag is set to true in the output, triggering the Python orchestrator to use embedding-based fallback.
- Prolog Exception: Caught by Python wrapper; returns a "Safe Default" decision (GPT-OSS, Medium Stakes, End-Stage Validation).

### 3.5 Routing Sequence

The logical execution flow within the route/2 predicate:

1. Normalize: Tokenize and clean input text.
2. Domain Classification: Match keywords to domain\_type.
3. Stakes Assessment: Calculate risk/complexity score based on domain and keywords.
4. Model Selection: Query model\_capability facts for the identified domain.
5. Validation Policy: Map stakes to validation\_policy.
6. Tool Detection: Check for multimodal triggers.
7. Confidence Scoring: Compute weighted average of domain, model, and stakes confidence.
8. Reasoning Generation: Construct a human-readable explanation string.